

# Class-23: JPA Entities & Repositories



by Pial Kanti Samadder

# Entity

- Entities in JPA are nothing but POJOs representing data that can be persisted in the database.
- An entity represents a table stored in a database.
- Every instance of an entity represents a row in the table.

## The *Entity* Annotation

Let's say we have a POJO called *Student*, which represents the data of a student, and we would like to store it in the database.

So let's define it by making use of the `@Entity` annotation. We must specify this annotation at the class level. **We must also ensure that the entity has a no-arg constructor and a primary key:**

```
@Entity  
public class Student {  
    // fields, getters and setters  
}
```

**Entity classes must not be declared *final*.**

# The *Id* Annotation

Each JPA entity must have a primary key that uniquely identifies it. The `@Id` annotation defines the primary key. We can generate the identifiers in different ways, which are specified by the `@GeneratedValue` annotation.

We can choose from four id generation strategies with the `strategy` element. **The value can be *AUTO*, *TABLE*, *SEQUENCE*, or *IDENTITY*:**

```
@Entity
public class Student {
    @Id
    @GeneratedValue(strategy=GenerationType.AUTO)
    private Long id;

    private String name;

    // getters and setters
}
```

If we specify `GenerationType.AUTO`, the JPA provider will use any strategy it wants to generate the identifiers.

# Primary Key Generation Strategies in JPA

JPA provides four strategies to generate primary keys for entities using the `@GeneratedValue` annotation.

## Syntax

```
@Id  
@GeneratedValue(strategy = GenerationType.XXX)  
private Long id;
```

# .GenerationType.AUTO

- This is the **default** strategy.
- JPA will **choose the strategy based on the underlying database** and the JPA provider (like Hibernate).
- Example: Hibernate may choose **SEQUENCE** for PostgreSQL, or **IDENTITY** for MySQL.

## Example

```
@Entity  
public class Product {  
    @Id  
    @GeneratedValue(strategy = GenerationType.AUTO)  
    private Long id;  
}
```

## Real-Time Use Case:

- Good for **portable applications** where the DBMS is not fixed.
- For example, in a **multi-tenant SaaS product** where customers may use different databases

# GenerationType.IDENTITY

- Uses the database's **auto-increment** feature.
- Primary key is generated by the DB at **insert time**.
- You can't batch insert with Hibernate using this strategy (limitation).

## Example (MySQL/PostgreSQL)

```
@Entity  
public class Customer {  
    @Id  
    @GeneratedValue(strategy = GenerationType.IDENTITY)  
    private Long id;  
}
```

## Real-Time Use Case:

- Used when your database supports auto-increment (e.g., **MySQL, PostgreSQL**).
- Common for **simple apps, CRUD apps**, or when inserting records one at a time.

## Note:

- Since ID is generated after insert, Hibernate can't predict it before saving, making batch inserts tricky.

# GenerationType.SEQUENCE

- Uses a **database sequence object**.
- Hibernate first fetches the next sequence value, then inserts the record.
- Allows **batch insert**.

## Example (PostgreSQL, Oracle):

```
@Entity  
public class Order {  
    @Id  
    @GeneratedValue(strategy = GenerationType.SEQUENCE, generator = "order_seq")  
    @SequenceGenerator(name = "order_seq", sequenceName = "order_sequence", allocationSize = 1)  
    private Long id;  
}
```

## Real-Time Use Case:

- Best for **enterprise systems** using **PostgreSQL or Oracle**, where performance and batch inserts are important.
- Example: A **banking system** or **e-commerce** app managing millions of transactions.

## Tip:

- You can increase `allocationSize` for fewer DB calls and better performance in batch operations

# GenerationType.TABLE

- Uses a **separate table** to simulate sequence behavior.
- Useful in **non-standard or older databases** that don't support **SEQUENCE** or **IDENTITY**.

## Example:

```
@Entity
public class Invoice {
    @Id
    @GeneratedValue(strategy = GenerationType.TABLE, generator = "invoice_gen")
    @TableGenerator(name = "invoice_gen", table = "id_gen", pkColumnName = "gen_name", valueColumnName = "gen_val",
    pkColumnValue = "invoice_id", allocationSize = 1)
    private Long id;
}
```

## Real-Time Use Case:

You're using an old or **non-standard database** that doesn't support sequences or auto-increment. You create a table called **id\_generator** that stores the current ID value for each entity.

# The *Table* Annotation

In most cases, **the name of the table in the database and the name of the entity won't be the same.**

In these cases, we can specify the table name using the *@Table* annotation:

```
@Entity  
@Table(name="students")  
public class Student {  
    // fields, getters and setters  
}
```

We can also mention the schema using the *schema* element:

```
@Entity  
@Table(name="students", schema="school")  
public class Student {  
    // fields, getters and setters  
}
```

Schema name helps to distinguish one set of tables from another.

If we don't use the *@Table* annotation, the name of the table will be the name of the entity.

# The *Column* Annotation

Just like the `@Table` annotation, we can use the `@Column` annotation to mention the details of a column in the table.

The `@Column` annotation has many elements such as `name`, `length`, `nullable`, and `unique`:

```
@Entity  
@Table(name="students")  
public class Student {  
    @Id  
    @GeneratedValue(strategy=GenerationType.AUTO)  
    private Long id;  
  
    @Column(name="student_name", length=50, nullable=false, unique=false)  
    private String name;  
  
    // other fields, getters and setters  
}
```

The `name` element specifies the name of the column in the table. The `length` element specifies its length. The `nullable` element specifies whether the column is nullable or not, and the `unique` element specifies whether the column is unique.

If we don't specify this annotation, the name of the column in the table will be the name of the field.

# The *Transient* Annotation

Sometimes, we may want to **make a field non-persistent**. We can use the `@Transient` annotation to do so. It specifies that the field won't be persisted.

For instance, we can calculate the age of a student from the date of birth.

So let's annotate the field `age` with the `@Transient` annotation:

```
@Entity  
@Table(name="students")  
public class Student {  
    @Id  
    @GeneratedValue(strategy=GenerationType.AUTO)  
    private Long id;  
  
    @Column(name="name", length=50, nullable=false)  
    private String name;  
  
    @Transient  
    private Integer age;  
  
    // other fields, getters and setters  
}
```

As a result, the field `age` won't be persisted in the table.

# What is a Transaction?

A **transaction** is a set of operations that must either **all succeed** or **all fail** as one unit.

## Example: Bank Transfer

- **Step 1:** Deduct money from Account A
- **Step 2:** Add money to Account B  
If one step fails, the **entire operation rolls back** — ensuring data consistency.

## In Spring Framework

A transaction groups database actions (e.g., `INSERT`, `UPDATE`, `DELETE`) so that:

-  All succeed → **Changes are committed**
-  Any fails → **All changes are rolled back**

## Goal

Ensure **data integrity**, **consistency**, and **reliability** in your application.

# ACID Properties of Transactions

Transactions usually use these four rules, known as **ACID**:

## A — Atomicity

- All operations in a transaction are **executed as one unit**.
- Either **everything succeeds or nothing happens**.
- On failure → all changes are **rolled back**.

## C — Consistency

- The database moves from **one valid state to another**.
- All **rules, constraints, and relationships** remain intact.
- Ensures **data integrity** after each transaction.

## I — Isolation

- Each transaction runs as if it's **the only one** in the system.
- Prevents **conflicts and interference** from concurrent transactions.
- Changes are **invisible to others** until committed.

## D — Durability

- Once committed, changes are **permanent**.
- Data **survives crashes or power failures**.
- The system **recovers to a consistent state** on restart.

# Why Do We Need Transactions?

Without transactions, multi-step or dependent operations can lead to **data inconsistency**.

## Example: Bank Transfer

If money is deducted from the sender's account but **not credited** to the receiver's — the system ends up with **corrupted data**.

## Transactions Ensure:

- **All-or-nothing execution** (maintains data integrity)
- **Automatic rollback** on failure
- **Consistent database state**
- **Controlled concurrency** through isolation levels
- **Flexible behavior** using propagation settings

Spring provides built-in transaction management features like **rollback rules**, **isolation levels**, and **propagation behaviors**, making it easier to maintain reliable and consistent data operations.

# Stages of a Transaction

## 1. Start

- Begins when a **transactional method** starts executing.
- All database operations within the method become part of **one transaction**.
- Changes are **temporary** until committed.

## 2. Commit

- Happens when the method **completes successfully** (no exceptions).
- All changes are **saved permanently** to the database.
- The transaction is now **complete and closed**.

## 3. Rollback

- Triggered when an exception occurs during execution.
- Undoes all changes made within the transaction.
- Returns the database to its original state.
- By default, Spring rolls back on unchecked (runtime) exceptions. You can customize this with the `rollbackFor` attribute.

# Transactions in Spring

The `@Transactional` annotation in Spring Framework simplifies working with transactions.

Instead of manually writing code to start, commit or rollback transactions, `@Transactional` handles all of that, and more, for us.

# How to Use @Transactional in Java Spring Methods

Using `@Transactional` is very simple in the Spring Framework, all you have to do is add the annotation to a method, most of the time it should be a service layer method. That essentially says Spring to make this method transactional.

Below is an example of using `@Transactional` in a simple mock example of using transactions for money transferring in the `BankService`:

```
public class BankService {  
    @Transactional  
    public void transferMoney(User sender, User recipient, double amount) {  
        removeFrom(sender, amount);  
        addTo(recipient, amount);  
    }  
  
    private void removeFrom(User user, double amount) {  
        // ... some logic for handling it would go here  
    }  
  
    private void addTo(User user, double amount) {  
        // ... some logic for handling it would go here  
    }  
}
```

In this example, we have two methods which will be encapsulated within an transaction.

- When the `transferMoney` method is called, `@Transactional` does all the work and transforms the operation into a transaction. The method is now a transaction in the start phase.
- If both the `removeFrom` and the `addTo` methods go through successfully without error, the transaction is committed. That is also known as commit stage.
- If there is an error in either of the methods, the transaction is rolled back to prevent partial updates. In case of an error, the transaction goes into the rollback stage.

# Considerations for @Transactional

## Best Practices

- **Apply on service layer methods** containing business logic, not repositories.
- Use `@Transactional(readOnly = true)` for **read-only operations**.
- Configure **propagation**, **isolation**, and **rollback rules** for advanced scenarios.

## Important Considerations

- **Private methods:** Transactions **do not apply** (pre-Spring 6) due to proxy limitations.
- **Internal method calls:** Calling a transactional method **within the same class bypasses the proxy**, so `@Transactional` is ignored.
- **Repositories:** Spring Data JPA auto-manages CRUD transactions, but **custom behavior still requires explicit annotation**.

## Key Takeaway

- **Transactional methods work only when called from outside the same class.**
- **Why:** Spring wraps your service in a “transaction manager” behind the scenes. Calls made **inside the same class** bypass this wrapper, so transactions are **not applied**.
- **Tip:** Always call transactional methods from another bean or from the client code, not from a private/internal method in the same class.

# What is Spring Data JPA Repository

**Goal:** Simplify database access by reducing boilerplate code.

**Key Idea:**

Spring Data JPA automatically provides CRUD (Create, Read, Update, Delete) operations without writing SQL or implementation code.

**Core Interface Hierarchy:**

```
Repository
└── CrudRepository
    └── PagingAndSortingRepository
        └── JpaRepository
```

**JpaRepository adds:**

- Pagination and sorting support
- Batch operations
- `Flush` and `saveAndFlush` methods

**Example:**

```
public interface UserRepository extends JpaRepository<User, Long> {  
}
```

# What Happens Under the Hood

## Spring Magic

- When your app starts, Spring creates a **proxy implementation** of the repository interface.
- This proxy handles:
  - Database connections (via EntityManager)
  - Transactions (auto-managed for CRUD)
  - Query generation based on method names

## Example:

```
userRepository.save(user); // Proxy executes SQL INSERT
```

## No manual JDBC or SQL needed!

## In summary

- JpaRepository is a high-level abstraction over JPA's EntityManager.
- Spring auto-generates implementations for your interfaces.
- Transactions are handled automatically for CRUD operations.

# Running Sample Project

- **Clone the Repository**

If you haven't cloned the repository yet, run the following command (ensure git is installed):

```
git clone https://github.com/PialKanti/Ostad-SpringBoot-Course.git
```

Then switch to the correct branch for today's class (replace with the actual branch name, e.g., class-23-jpa-entity):

```
git fetch  
git switch class-23-jpa-entity
```

Or,

If You Already Have the Repository Cloned, simply open your existing project folder and switch (or checkout) to the appropriate branch:

```
git fetch  
git switch class-23-jpa-entity
```

- **Set Up and Run PostgreSQL Database**

You can run PostgreSQL either via Docker or a desktop installation.

**Option 1: Run via Docker**

A compose.yml file is available in the root of the repository.

Run the following command from the project root:

```
docker compose up -d
```

This will start a PostgreSQL container automatically.

Or,

**Option 2: Run via PostgreSQL Desktop (Manual Setup)**

If you already have PostgreSQL installed locally:

1. Start your PostgreSQL server.
2. Create a new database named crud\_db if not exists.

- **Open the Project in IntelliJ IDEA**

1. Open IntelliJ IDEA.
2. Click **File → Open** and select the crud-sample folder inside the repository.
3. Let IntelliJ import Maven/Gradle dependencies automatically.